



Fraud Detection and Financial Risk Modeling

End-to-end fraud detection and financial risk modeling system (13M+ transactions) integrating EDA, feature engineering, LightGBM optimization, and explainable AI for business risk auditing.

Python 3.10+

LightGBM 4.x

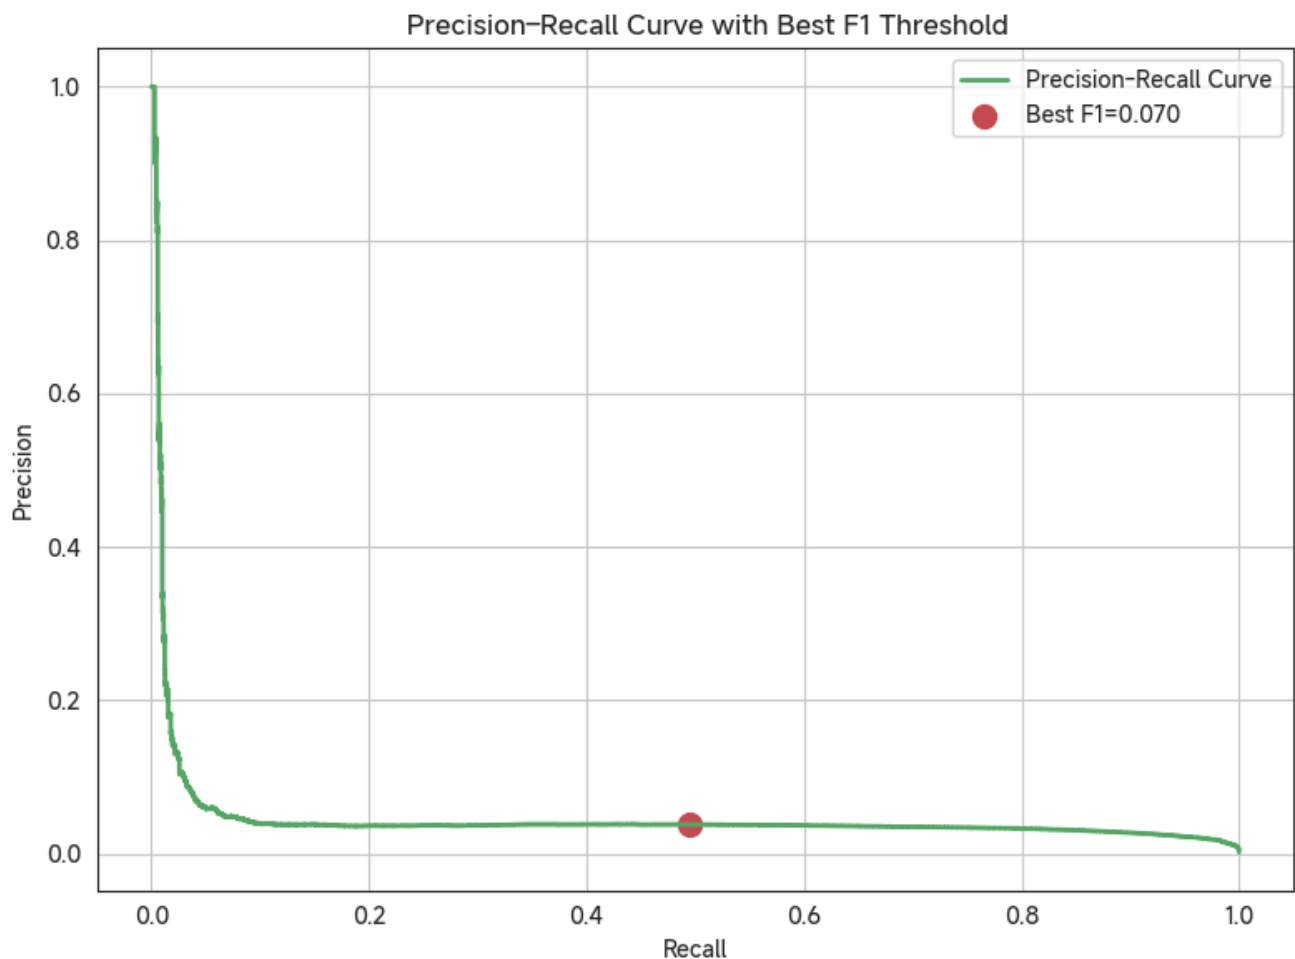
SHAP ExplainableAI


license MIT

Overview

This project builds a **data-driven fraud detection pipeline** capable of identifying suspicious financial transactions from over **13 million records**.

It combines **exploratory data analysis**, **feature engineering**, **imbalanced learning**, and **interpretable machine learning** to deliver actionable business insights.






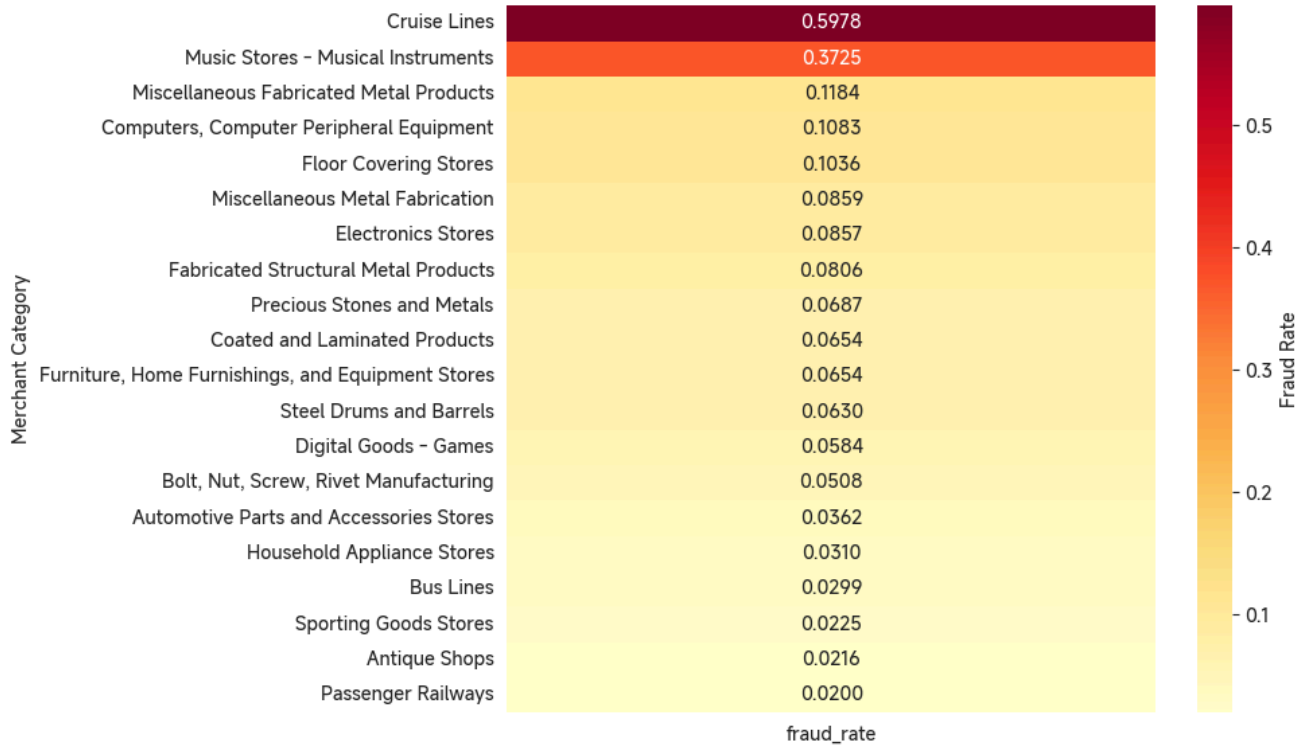
Workflow Architecture

Data Loading → EDA → Feature Engineering → Class Balancing →
LightGBM Training → Threshold Optimization → SHAP Interpretability →
Risk Report

Monthly Fraud Rate and Transaction Volume Over Time



Top 20 MCC Categories by Fraud Rate



Core Steps

Stage	Description
Data Preparation	Cleaned and merged <code>transactions</code> , <code>users</code> , <code>cards</code> , and <code>MCC</code> data (13M+ rows).
EDA & Visualization	Distribution, outlier, and geographic/industry-level fraud analysis.
Feature Engineering	Constructed 20+ numerical, categorical, and time-based features (<code>amount_log</code> , <code>is_refund</code> , <code>hour</code> , <code>mcc_desc</code> , etc.).
Class Balancing	Undersampling strategy (10× ratio) to stabilize LightGBM training.
Model Training	Tuned LightGBM (GBDT) with early stopping and AUC optimization.
Explainability (SHAP)	Interpreted feature contributions globally and locally (fraud driver visualization).
Threshold Tuning	Optimized precision–recall trade-off with decision curve visualization.
Business Application	Exported Top-200 high-risk transactions for manual verification.

Model Performance

Before & After Threshold Optimization

Metric	Default Threshold (0.5)	Tuned Threshold (0.1598)	Δ Improvement
AUC (Validation)	0.9718	0.9718	—
Precision	0.0374	0.0379	+0.0005
Recall	0.1601	0.4941	+0.3340
F1 Score	0.0606	0.0704	+0.0098
Detection Coverage	0.74%	2.28%	↑ Expanded audit scope

Interpretation:

- After tuning the decision threshold, Recall improved from **16.0% → 49.4%**, tripling fraud coverage while maintaining similar precision (~0.04).
- The optimized configuration balances model confidence and operational feasibility — ideal for financial audit pipelines where human verification capacity is limited.

Explainability

- Global Importance:** `amount_log` , `hour` , `mcc_desc` , and `client_mean` drive fraud prediction.
- Local Explanation:** SHAP force plots reveal how transaction timing and amount deviations trigger risk alerts.
- Interpretability Goal:** bridge **model confidence** and **financial analyst reasoning**.

Exported Deliverables

Output File	Description
<code>main.ipynb</code>	Full pipeline notebook with 19 modular cells
<code>main.html</code>	Rendered analysis report
<code>high_risk_transactions_top200.csv</code>	Top 200 suspicious transactions ranked by fraud probability
<code>.gitignore</code>	Excludes large data files (JSON/CSV over 100 MB)

Tech Stack

- **Python 3.10 + Pandas + NumPy**
- **LightGBM 4.x** (with callbacks for early stopping & AUC logging)
- **Matplotlib / Seaborn** (visual analytics)
- **SHAP 0.43.0** (explainable AI)
- **Scikit-learn 1.4+** (precision-recall & threshold tuning)

Business Impact

This project demonstrates a **scalable, interpretable fraud detection framework** ready for integration into financial risk control systems.

Author

Republic (GitHub: [Republic1024](#)) (AI & Data Science)

License

This project is released under the [MIT License](#).

